

Kubernetes Scenario Based Interview Questions with Expected Answers

Scenario 1: A Pod Keeps Crashing After Deployment

Question:

You've deployed a Pod, but it keeps crashing. How would you troubleshoot the issue?

Expected Answer:

- Check Pod status: `kubectl get pods`
 - Inspect logs: `kubectl logs <pod_name>`
 - Describe the Pod for events/errors: `kubectl describe pod <pod_name>`
 - Execute into the Pod to debug: `kubectl exec -it <pod_name> -- /bin/sh`
 - Validate environment variables and configurations
-

Scenario 2: High Traffic Causes Pod Failures

Question:

Your Kubernetes service is experiencing high traffic, and some requests are failing. How do you handle this?

Expected Answer:

- Scale the deployment: `kubectl scale deployment <name> --replicas=5`
 - Use HPA (Horizontal Pod Autoscaler): `kubectl autoscale deployment <name> --min=2 --max=10 --cpu-percent=80`
 - Ensure sufficient cluster resources using `kubectl top nodes`
 - Load balance traffic with an Ingress controller or Service
-

Scenario 3: A Service is Not Reachable

Question:

A deployed service is not accessible from outside the cluster. What could be wrong?

Expected Answer:

- Verify if the service type is ClusterIP, NodePort, or LoadBalancer
- Check the exposed ports using `kubectl get svc <service_name>`
- Ensure the service selector matches the pod labels

- Use `kubect1 port-forward` for debugging
-

Scenario 4: Persistent Volume Not Mounting Properly

Question:

You have configured a Persistent Volume (PV) and Persistent Volume Claim (PVC), but the Pod is unable to mount it. What do you do?

Expected Answer:

- Check PVC status: `kubect1 get pvc`
 - Verify if PV is correctly bound to PVC
 - Ensure the **storage class is supported** by the cluster
 - Describe the Pod and check for mount-related errors
-

Scenario 5: A Node Becomes Unavailable

Question:

One of the Kubernetes worker nodes suddenly becomes unavailable. What actions should you take?

Expected Answer:

- Check node status: `kubectl get nodes`
 - Inspect node conditions: `kubectl describe node <node_name>`
 - Investigate kubelet logs on the node
 - Drain and remove the node if necessary: `kubectl drain <node_name> --ignore-daemonsets --delete-local-data`
 - Ensure Pods are scheduled on healthy nodes
-

Scenario 6: Blue-Green Deployment in Kubernetes

Question:

How would you implement a Blue-Green Deployment in Kubernetes?

Expected Answer:

- Deploy a **new version (Green)** alongside the old one (**Blue**)
- Use **Service and Labels** to switch traffic
- Test the Green version, then update the service selector to point to it
- Roll back if needed using previous deployment history

Scenario 7: Securing Kubernetes Secrets

Question:

How can you securely manage sensitive credentials (API keys, passwords) in Kubernetes?

Expected Answer:

- Use Kubernetes **Secrets** instead of ConfigMaps
 - Store secrets in a secret management tool (e.g., HashiCorp Vault, AWS Secrets Manager)
 - Mount secrets as environment variables or files
 - Restrict access to Secrets using RBAC
-

Scenario 8: CI/CD Pipeline Deployment with Kubernetes

Question:

How would you automate deployments in Kubernetes using CI/CD?

Expected Answer:

- Use a **CI/CD tool** like Jenkins, GitHub Actions, or ArgoCD
- Build and push Docker images on code changes
- Apply manifests via `kubectl apply -f` or Helm
- Implement **Rolling Updates** for zero downtime

Scenario 9: A Rolling Update Fails Midway

Question:

You applied a Rolling Update, but some Pods failed, causing an inconsistent state. What's your approach to fixing this?

Expected Answer:

- Check rollout status: `kubectl rollout status deployment <deployment_name>`
 - View previous revisions: `kubectl rollout history deployment <deployment_name>`
 - Roll back if needed: `kubectl rollout undo deployment <deployment_name>`
 - Debug failing Pods: `kubectl describe pod <pod_name>`
-

Scenario 10: A Pod is Stuck in 'Pending' State

Question:

A Pod remains in the "Pending" state for too long. What do you investigate?

Expected Answer:

- Check `kubectl describe pod <pod_name>` for scheduling issues

- Verify node resources: `kubectl get nodes --show-labels`
 - Inspect Events: `kubectl get events --sort-by=.metadata.creationTimestamp`
 - If it's a volume issue, check `kubectl get pvc`
-

Scenario 11: Kubernetes Service is Not Routing Traffic Correctly

Question:

You deployed a service, but it's not distributing traffic to the Pods. What could be the issue?

Expected Answer:

- Check service selectors and Pod labels match
 - Ensure Pods are in "Running" state and have valid IPs
 - Validate service endpoints: `kubectl get endpoints <service_name>`
 - If using Ingress, check rules with `kubectl describe ingress`
-

Scenario 12: How Would You Handle a Multi-Cluster Deployment?

Question:

Your company needs to deploy Kubernetes workloads across multiple clusters. How would you approach it?

Expected Answer:

- Use **KubeFed (Kubernetes Federation)** for cross-cluster management
 - Leverage tools like **ArgoCD** or **FluxCD** for GitOps-based multi-cluster deployments
 - Use **Service Mesh (Istio/Linkerd)** to manage inter-cluster communication
-

Scenario 13: A Pod is Restarting Frequently

Question:

A Pod is constantly restarting. How do you troubleshoot?

Expected Answer:

- Check logs: `kubectl logs <pod_name> --previous`
- Describe the Pod: `kubectl describe pod <pod_name>`
- Investigate Liveness/Readiness probes in deployment YAML
- Verify resource limits (`kubectl get nodes` to check resource pressure)

Scenario 14: DNS Resolution is Not Working in Kubernetes

Question:

Your application Pods can't resolve DNS names of services inside the cluster. What do you do?

Expected Answer:

- Check CoreDNS status: `kubectl get pods -n kube-system | grep coredns`
 - Test DNS resolution: `kubectl exec -it <pod_name> -- nslookup <service_name>`
 - Verify service discovery settings in `/etc/resolv.conf` inside the Pod
 - Restart CoreDNS if necessary
-

Scenario 15: High Latency in Kubernetes Applications

Question:

Users report that your Kubernetes-based application is experiencing high latency. What's your debugging approach?

Expected Answer:

- Check resource usage with `kubectl top nodes` and `kubectl top pods`



Rakshita Belwal

- **Analyze network policies** (`kubect1 get networkpolicy`)
 - **Enable distributed tracing** with Jaeger/OpenTelemetry
 - **Check Ingress controller logs** for bottlenecks
-



Rakshita Belwal